

The Impact of Medicine to your Wallet

Today we will explore a data set dedicated to the cost of treatment of different patients. The cost of treatment depends on many factors: diagnosis, type of clinic, city of residence, age and so on. We have no data on the diagnosis of patients.

But we have other information that can help us to make a conclusion about the health of patients and practice regression analysis. In any case, I wish you to be healthy! Let's look at our data.

| Domain Knowledge

1. **Age:** Age is a significant factor in health insurance as older individuals are generally at a higher risk of health issues compared to younger individuals. Therefore, age often correlates with higher insurance charges.
2. **Sex:** Gender can play a role in health insurance costs. Some studies suggest that women tend to use more health care services, especially during childbearing years. However, the impact of gender on health insurance charges can vary depending on other factors.
3. **BMI (Body Mass Index):** BMI is a measure that attempts to quantify an individual's tissue mass. It's widely used as a simple method to assess whether a person has a healthy body weight for a given height. High BMI values are often associated with higher risk of chronic diseases like heart disease, diabetes, and certain types of cancer, which can lead to higher health care costs.
4. **Children:** The number of dependents a person has can influence health insurance costs. More dependents usually mean higher insurance charges because more individuals are covered under the same plan.
5. **Smoker:** Smoking is a significant risk factor for many health conditions, including heart disease, stroke, and cancer. Smokers often face higher health insurance charges due to these increased health risks.
6. **Region:** The region where the beneficiary lives can also affect health insurance charges. This could be due to variations in cost of living, health care costs, and health behaviors across different regions.
7. **Charges:** This is the individual medical costs billed by health insurance. It's the dependent variable in this dataset, and it's influenced by all the other factors. Understanding the relationships between the independent variables (age, sex, bmi, children, smoker, region) and charges is the main goal of this analysis.

| Import libraries

```
import pandas as pd
import numpy as np
import matplotlib

import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style("dark")
matplotlib.style.use('seaborn-notebook')
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = 'all'

# Suppress Scientific Notation in Pandas
pd.set_option('display.float_format', '{:.2f}'.format)

# for HD visualizations
%config InlineBackend.figure_format='retina'

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

# Evaluation
from sklearn import metrics
from sklearn.metrics import roc_curve
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error
from sklearn.metrics import recall_score, confusion_matrix,
precision_score, f1_score, accuracy_score, classification_report

C:\Users\prave\AppData\Local\Temp\ipykernel_16016\3204938556.py:9:
MatplotlibDeprecationWarning: The seaborn styles shipped by Matplotlib
are deprecated since 3.6, as they no longer correspond to the styles
shipped by seaborn. However, they will remain available as 'seaborn-
v0_8-<style>'. Alternatively, directly use the seaborn API instead.
  matplotlib.style.use('seaborn-notebook')

insurance_df = pd.read_csv(r"C:\Data\insurance.csv")

insurance_df.head(3)
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.90	0	yes	southwest	16884.92
1	18	male	33.77	1	no	southeast	1725.55
2	28	male	33.00	3	no	southeast	4449.46

```
insurance_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

!image-2.png

| Handling Duplicates

```
# Handle duplicates
duplicate_rows_data = insurance_df[insurance_df.duplicated()]
print("number of duplicate rows: ", duplicate_rows_data.shape)

number of duplicate rows:  (1, 7)

insurance_df[insurance_df.duplicated()]

   age  sex  bmi  children  smoker  region  charges
581  19  male  30.59         0     no  northwest  1639.56

insurance_df = insurance_df.drop_duplicates()
```

| Uniqueness

```
# Loop through each column and count the number of distinct values
for column in insurance_df.columns:
    num_distinct_values = len(insurance_df[column].unique())
    print(f"{column}: {num_distinct_values} distinct values")

age: 47 distinct values
sex: 2 distinct values
bmi: 548 distinct values
children: 6 distinct values
smoker: 2 distinct values
region: 4 distinct values
charges: 1337 distinct values
```

| Missing Values

```
# Missing or Null values
insurance_df.isna().sum()

age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

| Describe the Data


```
insurance_df.describe().style.format("{:.2f}")

<pandas.io.formats.style.Styler at 0x1647b55cf40>
```


image.png

Univariate analysis


```
# Histogram of age distribution
plt.figure(figsize=(8, 5))
plt.hist(insurance_df['age'], bins=20, edgecolor='black')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.title('Age Distribution')
plt.show();
```




```
# Bar plot of gender distribution
plt.figure(figsize=(6,5))
sns.countplot(data = insurance_df, x = 'sex')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.title('Gender Distribution')
plt.show();
```



```
# Histogram of BMI values  
plt.figure(figsize=(8, 4))  
plt.hist(insurance_df['bmi'], bins=20, edgecolor='black')  
plt.xlabel('BMI')  
plt.ylabel('Frequency')  
plt.title('BMI Distribution')  
plt.show();
```




```
# Visualization : Bar plot of the number of children/dependents
plt.figure(figsize=(8,4))
sns.countplot(data = insurance_df, x='children')
plt.xlabel('Number of Children')
plt.ylabel('Count')
plt.title('Number of Children Distribution')
plt.show();
```

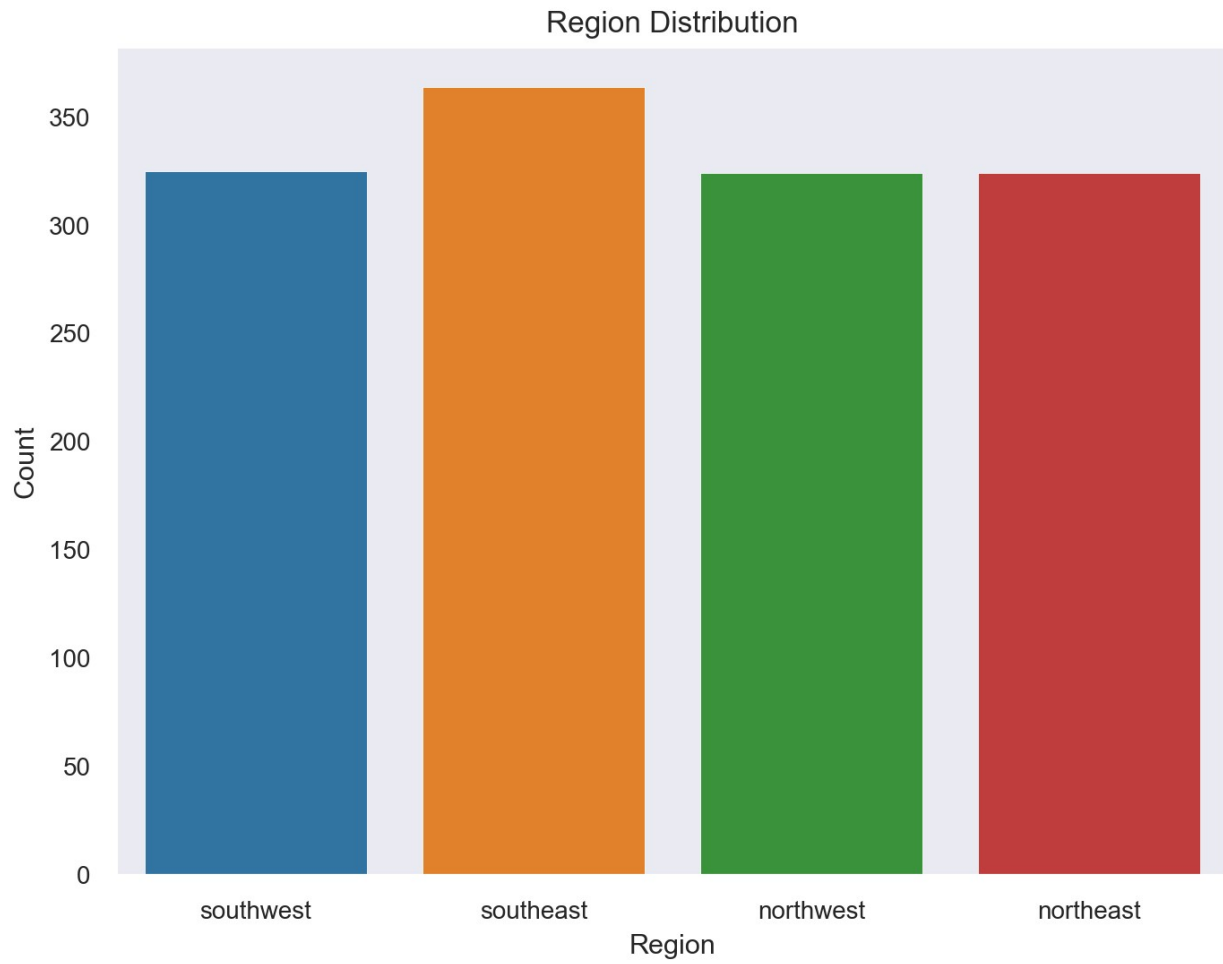


```
# Pie chart of smoker vs. non-smoker distribution
plt.figure(figsize=(5,5))
insurance_df['smoker'].value_counts().plot(kind='pie', autopct='%1.1f%%')
plt.title('Smoker vs. Non-Smoker Distribution')
plt.show();
```

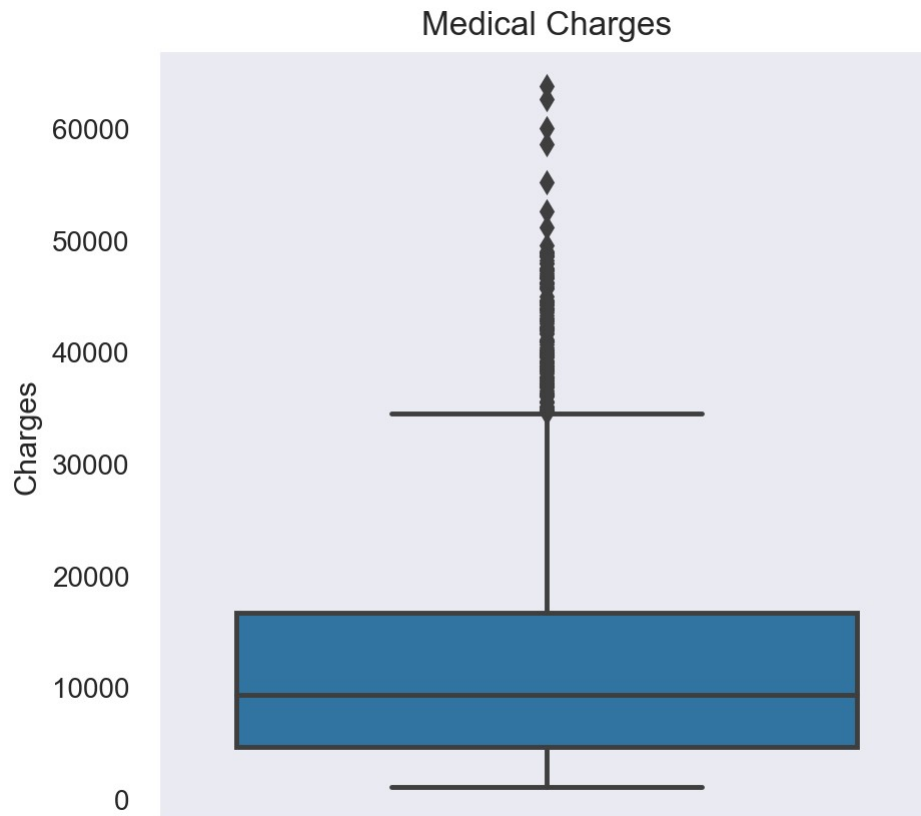

Smoker vs. Non-Smoker Distribution



```
# Bar plot of region distribution
plt.figure(figsize=(8, 6))
sns.countplot(data = insurance_df, x='region')
plt.xlabel('Region')
plt.ylabel('Count')
plt.title('Region Distribution')
plt.show();
```




```
# Box plot of medical charges  
plt.figure(figsize=(5,5))  
sns.boxplot(data = insurance_df, y='charges')  
plt.ylabel('Charges')  
plt.title('Medical Charges')  
plt.show();
```

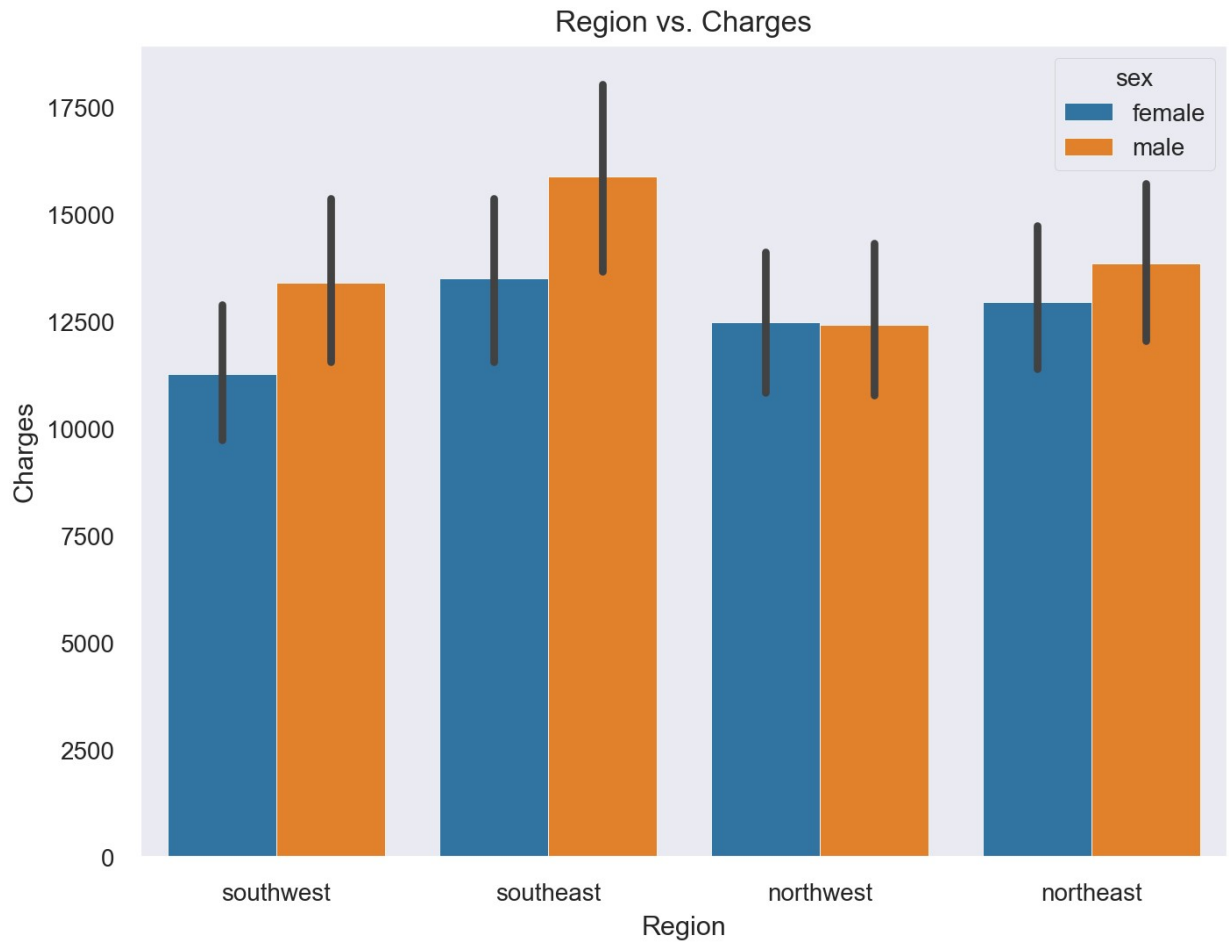


Bi-Variate Analysis


```
# Scatter plot of age vs. charges  
plt.figure(figsize=(8,5))  
plt.scatter(insurance_df['age'], insurance_df['charges'])  
plt.xlabel('Age')  
plt.ylabel('Charges')  
plt.title('Age vs. Charges')  
plt.show();
```



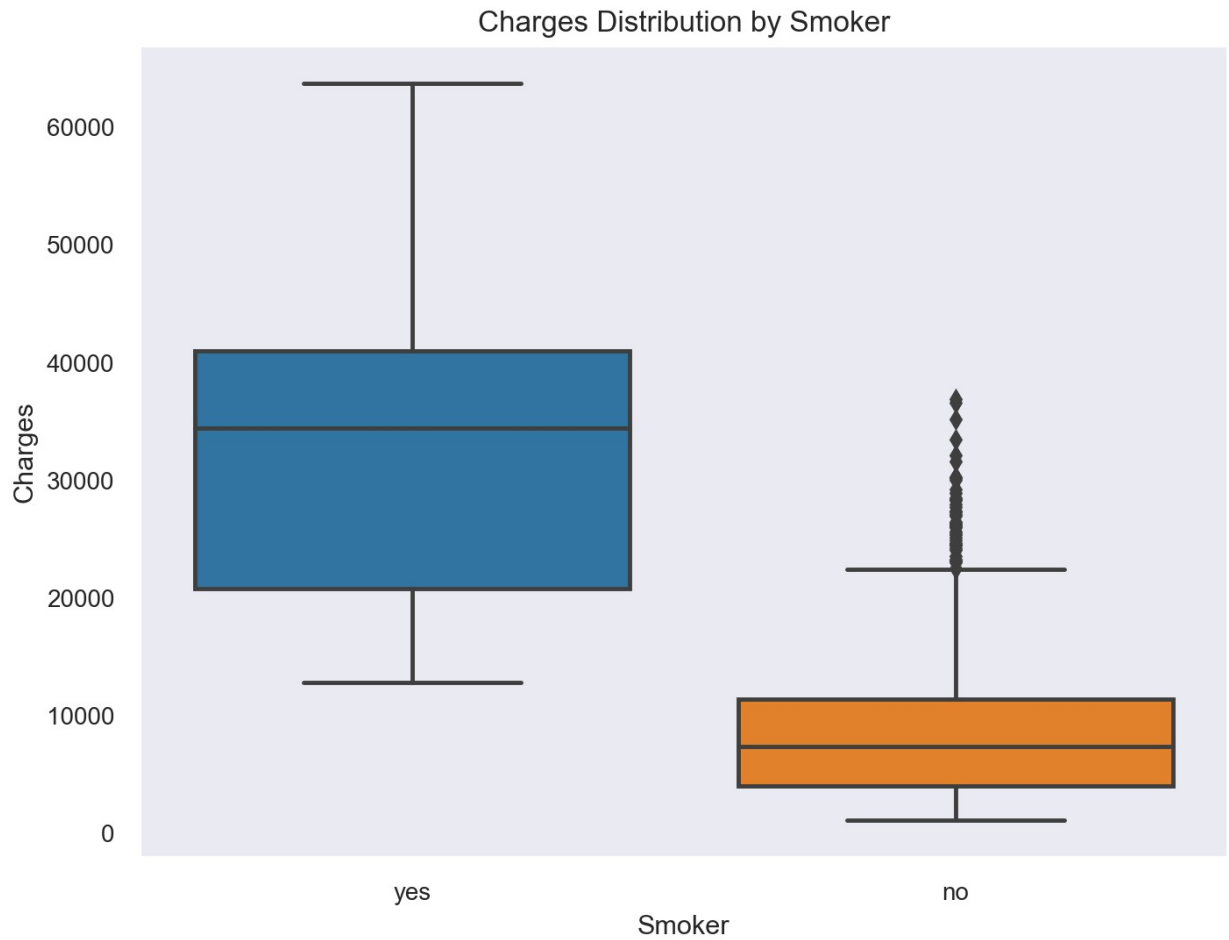
```
# Stacked bar plot of region vs. charges
plt.figure(figsize=(8, 6))
sns.barplot(data = insurance_df, x='region', y='charges', hue='sex')
plt.xlabel('Region')
plt.ylabel('Charges')
plt.title('Region vs. Charges')
plt.show();
```




```
# Box plot of charges grouped by gender
plt.figure(figsize=(8, 6))
sns.boxplot(data=insurance_df, x='sex', y='charges')
plt.xlabel('Gender')
plt.ylabel('Charges')
plt.title('Charges Distribution by Gender')
plt.show();
```



```
# Box plot of charges grouped by smoker vs. non-smoker  
plt.figure(figsize=(8, 6))  
sns.boxplot(data=insurance_df, x='smoker', y='charges')  
plt.xlabel('Smoker')  
plt.ylabel('Charges')  
plt.title('Charges Distribution by Smoker')  
plt.show();
```





```
# Box plot of charges grouped by the number of children  
plt.figure(figsize=(8, 6))  
sns.boxplot(data=insurance_df, x='children', y='charges')  
plt.xlabel('Number of Children')  
plt.ylabel('Charges')  
plt.title('Charges Distribution by Number of Children')  
plt.show();
```



```
plt.figure(figsize=(15,5))

plt.subplot(121)
sns.boxplot(insurance_df.charges)

plt.subplot(122)
sns.histplot(insurance_df.charges.astype(float), kde = True)
plt.axvline(insurance_df.charges.astype(float).mean(), color = 'blue')
plt.axvline(insurance_df.charges.astype(float).median(), color = 'red');
```

```
insurance_df.groupby('children').agg(['mean', 'min', 'max'])['charges']
```

	mean	min	max
children			
0	12384.70	1121.87	63770.43
1	12731.17	1711.03	58571.07
2	15073.56	2304.00	49577.66
3	15355.32	3443.06	60021.40
4	13850.66	4504.66	40182.25
5	8786.04	4687.80	19023.26

Outliers

```
plt.figure(figsize=(15,3))
sns.heatmap(insurance_df.isnull(),cbar=False,cmap='summer',yticklabels=False)
plt.title('Missing value in the dataset');
```



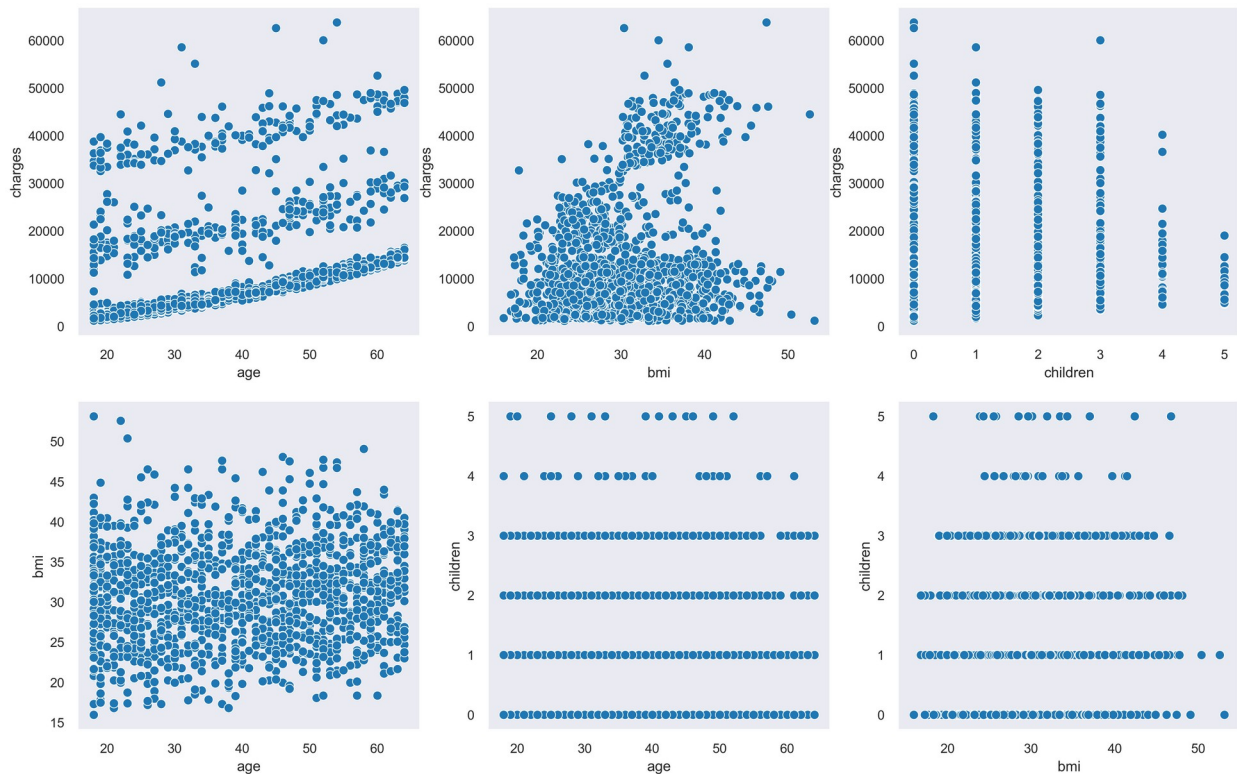
There is no missing values in data set

```
plt.figure(figsize=(16,10), dpi = 200)
plt.subplot(231)
sns.scatterplot(data = insurance_df, x = 'age', y = 'charges')
plt.subplot(232)
sns.scatterplot(data =insurance_df, x = 'bmi', y = 'charges')
plt.subplot(233)
```

```

sns.scatterplot(data =insurance_df, x = 'children', y = 'charges')
plt.subplot(234)
sns.scatterplot(data =insurance_df, x = 'age', y = 'bmi')
plt.subplot(235)
sns.scatterplot(data =insurance_df, x = 'age', y = 'children')
plt.subplot(236)
sns.scatterplot(data =insurance_df, x = 'bmi', y = 'children');

```



! image.png

Data Preprocessing

| Encoding

Machine learning algorithms cannot work with categorical data directly, categorical data must be converted to number.

| Label Encoding

Label encoding refers to transforming the word labels into numerical form so that the algorithms can understand how to operate on them.

| One Hot encoding

A One hot encoding is a representation of categorical variable as binary vectors. It allows the representation of categorical data to be more expressive. This first requires that the categorical

values be mapped to integer values, that is label encoding. Then, each integer value is represented as a binary vector that is all zero values except the index of the integer, which is marked with a 1.

```
X = insurance_df[['age', 'sex', 'bmi', 'smoker', 'children']]
y = insurance_df['charges']

X.head()
y.head()
```

	age	sex	bmi	smoker	children
0	19	female	27.90	yes	0
1	18	male	33.77	no	1
2	28	male	33.00	no	3
3	33	male	22.70	no	0
4	32	male	28.88	no	0

0	16884.92
1	1725.55
2	4449.46
3	21984.47
4	3866.86

Name: charges, dtype: float64

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.25, random_state = 100)

print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```

(1002, 5) (1002,)
(335, 5) (335,)

Seperating Categorical and numerical

```
X_train.head()
```

	age	sex	bmi	smoker	children
1000	30	male	22.99	yes	2
885	32	male	28.93	yes	1
517	45	male	30.50	no	2
1132	57	male	40.28	no	0
1188	43	female	25.27	yes	1

```
X_train_cat = X_train.select_dtypes(include=['object'])
X_train_cat.head()
```

	sex	smoker
1000	male	yes

885	male	yes
517	male	no
1132	male	no
1188	female	yes

```
X_train_num = X_train.select_dtypes(include=['int64', 'float64'])
```

```
X_train_num.head()
```

	age	bmi	children
1000	30	22.99	2
885	32	28.93	1
517	45	30.50	2
1132	57	40.28	0
1188	43	25.27	1

Scaling the Numerical features

```
X_train_num.head()
```

	age	bmi	children
1000	30	22.99	2
885	32	28.93	1
517	45	30.50	2
1132	57	40.28	0
1188	43	25.27	1

```
# scaling the numerical features
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
# column names are (annoyingly) lost after Scaling
```

```
# (i.e. the dataframe is converted to a numpy ndarray)
```

```
X_train_num_rescaled = pd.DataFrame(scaler.fit_transform(X_train_num),
                                     columns = X_train_num.columns,
                                     index = X_train_num.index)
```

```
X_train_num_rescaled.head()
```

	age	bmi	children
1000	-0.64	-1.23	0.79
885	-0.50	-0.27	-0.07
517	0.42	-0.01	0.79
1132	1.27	1.58	-0.92
1188	0.28	-0.86	-0.07

```
X_train_num.describe()
```

	age	bmi	children
count	1002.00	1002.00	1002.00
mean	39.08	30.57	1.08
std	14.10	6.14	1.18
min	18.00	15.96	0.00
25%	26.00	26.08	0.00
50%	39.00	30.30	1.00
75%	51.00	34.80	2.00
max	64.00	53.13	5.00

```
print("Number of Numerical Features:", scaler.n_features_in_)
print("Mean of each column:", scaler.mean_)
print("Std of each column:", np.sqrt(scaler.var_))
```

```
Number of Numerical Features: 3
Mean of each column: [39.08283433 30.57117265  1.07684631]
Std of each column: [14.09138002  6.14168693  1.17572029]
```

Applying OneHotEncoding on Categorical Columns

```
X_train_cat['sex'].value_counts(normalize=True)
```

```
male    0.51
female  0.49
Name: sex, dtype: float64
```

```
X_train_cat['smoker'].value_counts(normalize=True)
```

```
no    0.79
yes   0.21
Name: smoker, dtype: float64
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
encoder = OneHotEncoder(drop='first', sparse=False)
```

```
# column names are (annoyingly) lost after OneHotEncoding
# (i.e. the dataframe is converted to a numpy ndarray)
```

```
X_train_cat_ohe = pd.DataFrame(encoder.fit_transform(X_train_cat),
                                columns=encoder.get_feature_names_out(X_train_cat.columns),
                                index = X_train_cat.index)
```

```
X_train_cat_ohe.head()
```

	sex_male	smoker_yes
1000	1.00	1.00
885	1.00	1.00
517	1.00	0.00

1132	1.00	0.00
1188	0.00	1.00

Concatinating the Encoded Categorical Features and Rescaled Numerical Features

```
X_train_transformed = pd.concat([X_train_num_rescaled,
X_train_cat_ohe], axis=1)
```

```
X_train_transformed.head()
```

	age	bmi	children	sex_male	smoker_yes
1000	-0.64	-1.23	0.79	1.00	1.00
885	-0.50	-0.27	-0.07	1.00	1.00
517	0.42	-0.01	0.79	1.00	0.00
1132	1.27	1.58	-0.92	1.00	0.00
1188	0.28	-0.86	-0.07	0.00	1.00

Preparing Test Data

```
X_test.head()
```

	age	sex	bmi	smoker	children
12	23	male	34.40	no	0
306	28	female	27.50	no	2
318	44	female	27.64	no	0
816	24	female	24.23	no	0
157	18	male	25.18	yes	0

```
X_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 335 entries, 12 to 625
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         335 non-null    int64
1   sex         335 non-null    object
2   bmi         335 non-null    float64
3   smoker      335 non-null    object
4   children    335 non-null    int64
dtypes: float64(1), int64(2), object(2)
memory usage: 15.7+ KB
```

```
X_test_cat = X_test.select_dtypes(include=['object'])
```

```
X_test_cat.head()
```

```
X_test_num = X_test.select_dtypes(include=['int64', 'float64'])
```

```
X_test_num.head()
```

	sex	smoker
12	male	no
306	female	no
318	female	no
816	female	no
157	male	yes

	age	bmi	children
12	23	34.40	0
306	28	27.50	2
318	44	27.64	0
816	24	24.23	0
157	18	25.18	0

```
X_test_cat_ohe = pd.DataFrame(encoder.transform(X_test_cat),  
                                columns=encoder.get_feature_names_out(X_test_cat.columns),  
                                index = X_test_cat.index)
```

```
X_test_cat_ohe.head()
```

	sex_male	smoker_yes
12	1.00	0.00
306	0.00	0.00
318	0.00	0.00
816	0.00	0.00
157	1.00	1.00

```
X_test_num_rescaled = pd.DataFrame(scaler.transform(X_test_num),  
                                    columns = X_test_num.columns,  
                                    index = X_test_num.index)
```

```
X_test_num_rescaled.head()
```

	age	bmi	children
12	-1.14	0.62	-0.92
306	-0.79	-0.50	0.79
318	0.35	-0.48	-0.92
816	-1.07	-1.03	-0.92
157	-1.50	-0.88	-0.92

```
X_test_transformed = pd.concat([X_test_num_rescaled, X_test_cat_ohe],  
                                axis=1)
```

```
X_test_transformed.head()
```

	age	bmi	children	sex_male	smoker_yes
12	-1.14	0.62	-0.92	1.00	0.00

306	-0.79	-0.50	0.79	0.00	0.00
318	0.35	-0.48	-0.92	0.00	0.00
816	-1.07	-1.03	-0.92	0.00	0.00
157	-1.50	-0.88	-0.92	1.00	1.00

Variance inflation factor

```
from patsy import dmatrices
from statsmodels.stats.outliers_influence import
variance_inflation_factor

vif = pd.DataFrame()
vif['VIF'] = [variance_inflation_factor(X_train_transformed.values, i)
for i in range(X_train_transformed.shape[1])]
vif['variable'] = X_train_transformed.columns
```

vif

	VIF	variable
0	1.01	age
1	1.01	bmi
2	1.00	children
3	1.16	sex_male
4	1.15	smoker_yes

```
vif = pd.DataFrame()
vif['VIF'] = [variance_inflation_factor(X_test_transformed.values, i)
for i in range(X_test_transformed.shape[1])]
vif['variable'] = X_test_transformed.columns
```

vif

	VIF	variable
0	1.04	age
1	1.04	bmi
2	1.01	children
3	1.19	sex_male
4	1.19	smoker_yes

| Model Building

Linear Regression

```
from sklearn.linear_model import LinearRegression
LR_regressor = LinearRegression()
LR_regressor.fit(X_train_transformed, y_train)
```



```

LinearRegression()

y_test_pred = LR_regressor.predict(X_test_transformed)

#Testing the model accuracy

r2 = r2_score(y_test, y_test_pred) # Getting the R-squared (R²)
mae = mean_absolute_error(y_test, y_test_pred) # Getting the Mean Absolute Error
mse = mean_squared_error(y_test, y_test_pred) # Getting the Mean Squared Error
rmse = np.sqrt(mse) # Getting the Root Mean Squared Error

#Defining a DataFrame
results = pd.DataFrame([['Linear Regression', r2*100, mae, mse, rmse]],
                        columns=['Model', 'R-squared (R²)', 'Mean Absolute Error', 'Mean Squared Error', 'Root Mean Squared Error'])

#Printing the Model Results with different Metrics
results

```

	Model	R-squared (R²)	Mean Absolute Error	Mean Squared Error	Root Mean Squared Error
0	Linear Regression	75.40	3938.02	31503873.50	5612.83

```

print('Linear Regression Train Score is : ' ,
LR_regressor.score(X_train_transformed, y_train))
print('Linear Regression Test Score is : ' ,
LR_regressor.score(X_test_transformed, y_test))

Linear Regression Train Score is : 0.747701374571325
Linear Regression Test Score is : 0.7540307970555084

```

Residuals (Difference between y_test and y_pred)

```

test_residual= y_test - y_test_pred

medical_LR = pd.DataFrame({'Actual': y_test, 'Predicted': y_test_pred,
'diff':(abs(y_test-y_test_pred))})

medical_LR.head()

```

	Actual	Predicted	diff
12	1826.84	4798.71	2971.87
306	20177.67	5005.99	15171.68
318	7421.19	8305.25	884.06
816	2842.76	1823.34	1019.42
157	15518.18	24129.98	8611.80

KNN Regression

```

from sklearn.neighbors import KNeighborsRegressor
KNN_regressor = KNeighborsRegressor()
KNN_regressor.fit(X_train_transformed, y_train)

KNeighborsRegressor()

y_test_pred = KNN_regressor.predict(X_test_transformed)

#Testing the model accuracy

r2 = r2_score(y_test, y_test_pred) # Getting the R-squared (R²)
mae = mean_absolute_error(y_test, y_test_pred) # Getting the Mean Absolute Error
mse = mean_squared_error(y_test, y_test_pred) # Getting the Mean Squared Error
rmse = np.sqrt(mse) # Getting the Root Mean Squared Error

#Defining a DataFrame
knn_results = pd.DataFrame([['KNN Regression', r2*100, mae, mse, rmse]],
                            columns=['Model', 'R-squared (R²)', 'Mean Absolute Error', 'Mean Squared Error', 'Root Mean Squared Error'])

#Printing the Model Results with different Metrics
results = results.append(knn_results , ignore_index=True)
results

```

	Model	R-squared (R²)	Mean Absolute Error	Mean Squared Error	Root Mean Squared Error
0	Linear Regression	75.40	3938.02	31503873.50	5612.83
1	KNN Regression	79.97	3006.35	25660048.66	5065.57

```
print('KNN Regression Train Score is : ' ,
KNN_regressor.score(X_train_transformed, y_train))
print('KNN Regression Test Score is : ' ,
KNN_regressor.score(X_test_transformed, y_test))

KNN Regression Train Score is : 0.8386114862600965
KNN Regression Test Score is : 0.7996569622824252
```

Residuals (Difference between y_test and y_pred)

```
test_residual= y_test - y_test_pred

medical_knn = pd.DataFrame({'Actual': y_test, 'Predicted':
y_test_pred, 'diff':(abs(y_test-y_test_pred))})

medical_knn.head()
```

	Actual	Predicted	diff
12	1826.84	1652.96	173.88
306	20177.67	4587.75	15589.92
318	7421.19	6863.31	557.89
816	2842.76	7573.04	4730.28
157	15518.18	19947.37	4429.19

Decision Tree Regression

```
from sklearn.tree import DecisionTreeRegressor

DT_regressor = DecisionTreeRegressor()
DT_regressor.fit(X_train_transformed, y_train)

DecisionTreeRegressor()

y_test_pred = DT_regressor.predict(X_test_transformed)

#Testing the model accuracy

r2 = r2_score(y_test, y_test_pred) # Getting the R-
squared (R²)
mae = mean_absolute_error(y_test, y_test_pred) # Getting the Mean
Absolute Error
mse = mean_squared_error(y_test, y_test_pred) # Getting the Mean
Squared Error
rmse = np.sqrt(mse) # Getting the Root
Mean Squared Error

#Defining a DataFrame
dt_results = pd.DataFrame([['Decison-tree Regression', r2*100, mae,
mse, rmse]],
```

```
columns=['Model', 'R-squared (R²)', 'Mean Absolute Error', 'Mean Squared Error', 'Root Mean Squared Error']
```

#Printing the Model Results with different Metrics

```
results = results.append(dt_results , ignore_index=True)
results
```

	Model	R-squared (R²)	Mean Absolute Error \
0	Linear Regression	75.40	3938.02
1	KNN Regression	79.97	3006.35
2	Decison-tree Regression	58.68	3648.23

	Mean Squared Error	Root Mean Squared Error
0	31503873.50	5612.83
1	25660048.66	5065.57
2	52924035.93	7274.89

```
print('Decision Regression Train Score is : ' ,
DT_regressor.score(X_train_transformed, y_train))
print('Decision Regression Test Score is : ' ,
DT_regressor.score(X_test_transformed, y_test))
```

```
Decision Regression Train Score is : 0.9988021709249079
Decision Regression Test Score is : 0.5867910358585947
```

Residuals (Difference between y_test and y_pred)

```
test_residual= y_test - y_test_pred
```

```
medical_DT = pd.DataFrame({'Actual': y_test, 'Predicted': y_test_pred,
'diff':(abs(y_test-y_test_pred))})
```

```
medical_DT.head()
```

	Actual	Predicted	diff
12	1826.84	1837.28	10.44
306	20177.67	4340.44	15837.23
318	7421.19	7626.99	205.80
816	2842.76	22395.74	19552.98
157	15518.18	16232.85	714.67

Hyperparametric Tuning for Decision Tree

```
from sklearn.model_selection import RandomizedSearchCV
```

```
parameters={"splitter":["best","random"],
            "max_depth" : [1,3,5,7,9],
            "min_samples_leaf":[1,2,3,4,5,],
            "max_features":["auto","sqrt"],
            "max_leaf_nodes":[5,10,15] }
```

```

hyper_tuning = RandomizedSearchCV(estimator=DT_regressor,
param_distributions = parameters,
                                cv = 2, n_iter = 10, n_jobs=-1)

hyper_tuning.fit(X_train_transformed, y_train)

RandomizedSearchCV(cv=2, estimator=DecisionTreeRegressor(), n_jobs=-1,
                    param_distributions={'max_depth': [1, 3, 5, 7, 9],
                                        'max_features': ['auto',
'sqrt'],
                                        'max_leaf_nodes': [5, 10, 15],
                                        'min_samples_leaf': [1, 2, 3,
4, 5],
                                        'splitter': ['best',
'random']})

hyper_tuning.best_params_    ## getting the best parameters
{'splitter': 'best',
 'min_samples_leaf': 5,
 'max_leaf_nodes': 5,
 'max_features': 'auto',
 'max_depth': 5}

DT_hyper =
DecisionTreeRegressor(splitter='best',min_samples_leaf=4,max_leaf_node
s=15,max_features='auto',max_depth=7)

DT_hyper.fit(X_train_transformed,y_train)

DecisionTreeRegressor(max_depth=7, max_features='auto',
max_leaf_nodes=15,
                    min_samples_leaf=4)

tuning_predict = DT_hyper.predict(X_test_transformed)

print('Decision Tree Regressor hyperparameter score
training :',hyper_tuning.score(X_train_transformed,y_train))

Decision Tree Regressor hyperparameter score training :
0.8361318429771127

print('Decision Tree Regressor hyperparameter score testing :',
hyper_tuning.score(X_test_transformed,y_test))

Decision Tree Regressor hyperparameter score testing :
0.8372962465579079

```

Random Forest Regression

```
from sklearn.ensemble import RandomForestRegressor
RF_regressor = RandomForestRegressor()
RF_regressor.fit(X_train_transformed, y_train)

RandomForestRegressor()

y_test_pred = RF_regressor.predict(X_test_transformed)

#Testing the model accuracy

r2 = r2_score(y_test, y_test_pred)           # Getting the R-squared (R²)
mae = mean_absolute_error(y_test, y_test_pred) # Getting the Mean Absolute Error
mse = mean_squared_error(y_test, y_test_pred) # Getting the Mean Squared Error
rmse = np.sqrt(mse)                          # Getting the Root Mean Squared Error

#Defining a DataFrame
RF_results = pd.DataFrame([['Random Forest Regression', r2*100, mae, mse, rmse]],
                           columns=['Model', 'R-squared (R²)', 'Mean Absolute Error', 'Mean Squared Error', 'Root Mean Squared Error'])

#Printing the Model Results with different Metrics
results = results.append(RF_results, ignore_index=True)
results
```

	Model	R-squared (R²)	Mean Absolute Error \
0	Linear Regression	75.40	3938.02
1	KNN Regression	79.97	3006.35
2	Decison-tree Regression	58.68	3648.23
3	Random Forest Regression	82.55	2701.47

	Mean Squared Error	Root Mean Squared Error
0	31503873.50	5612.83
1	25660048.66	5065.57
2	52924035.93	7274.89
3	22344392.41	4726.99

```
print('Random Forest Regression Train Score is : ',
      RF_regressor.score(X_train_transformed, y_train))
print('Random Forest Regression Test Score is : ',
      RF_regressor.score(X_test_transformed, y_test))
```

```
Random Forest Regression Train Score is : 0.9750853574499634
Random Forest Regression Test Score is : 0.8255442337328778
```

Residuals (Difference between y_test and y_pred)

```
test_residual= y_test - y_test_pred

medical_RF = pd.DataFrame({'Actual': y_test, 'Predicted': y_test_pred,
'diff':(abs(y_test-y_test_pred))})

medical_RF.head()
```

	Actual	Predicted	diff
12	1826.84	2768.14	941.30
306	20177.67	4594.38	15583.29
318	7421.19	16963.18	9541.99
816	2842.76	11651.18	8808.42
157	15518.18	16475.83	957.65

Hyperparametric Tuning for Random Forest

```
grid={'n_estimators':[int(x) for x in
np.linspace(start=100,stop=1200,num=12)],
      'max_features':['auto','log2'],
      'max_depth':[int(x) for x in np.linspace(5,30,num=6)],
      'min_samples_split':[2,5,10,15,100],
      'min_samples_leaf':[1,2,5,10]}

hyper_tuning=RandomizedSearchCV(estimator=RF_regressor,param_distribut
ions=grid,n_iter=10,cv=5,verbose=5,random_state=2)

hyper_tuning.fit(X_train_transformed,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[CV 1/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=800;; score=0.855 total time=3.2s

[CV 2/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=800;; score=0.878 total time=3.1s

[CV 3/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=800;; score=0.768 total time=3.0s

[CV 4/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=800;; score=0.824 total time=3.0s

[CV 5/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=800;; score=0.847 total time=3.0s

[CV 1/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=5, n_estimators=1000;; score=0.877 total time=4.8s

[CV 2/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=5, n_estimators=1000;; score=0.888 total time=

4.8s
[CV 3/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=5, n_estimators=1000;; score=0.778 total time= 5.3s
[CV 4/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=5, n_estimators=1000;; score=0.841 total time= 5.0s
[CV 5/5] END max_depth=30, max_features=auto, min_samples_leaf=5, min_samples_split=5, n_estimators=1000;; score=0.854 total time= 5.0s
[CV 1/5] END max_depth=20, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=100;; score=0.853 total time= 0.3s
[CV 2/5] END max_depth=20, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=100;; score=0.880 total time= 0.3s
[CV 3/5] END max_depth=20, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=100;; score=0.767 total time= 0.3s
[CV 4/5] END max_depth=20, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=100;; score=0.823 total time= 0.3s
[CV 5/5] END max_depth=20, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=100;; score=0.848 total time= 0.3s
[CV 1/5] END max_depth=25, max_features=log2, min_samples_leaf=10, min_samples_split=2, n_estimators=800;; score=0.856 total time= 2.6s
[CV 2/5] END max_depth=25, max_features=log2, min_samples_leaf=10, min_samples_split=2, n_estimators=800;; score=0.877 total time= 2.7s
[CV 3/5] END max_depth=25, max_features=log2, min_samples_leaf=10, min_samples_split=2, n_estimators=800;; score=0.753 total time= 2.8s
[CV 4/5] END max_depth=25, max_features=log2, min_samples_leaf=10, min_samples_split=2, n_estimators=800;; score=0.822 total time= 2.7s
[CV 5/5] END max_depth=25, max_features=log2, min_samples_leaf=10, min_samples_split=2, n_estimators=800;; score=0.819 total time= 2.9s
[CV 1/5] END max_depth=30, max_features=auto, min_samples_leaf=2, min_samples_split=100, n_estimators=700;; score=0.854 total time= 2.8s
[CV 2/5] END max_depth=30, max_features=auto, min_samples_leaf=2, min_samples_split=100, n_estimators=700;; score=0.877 total time= 2.6s
[CV 3/5] END max_depth=30, max_features=auto, min_samples_leaf=2, min_samples_split=100, n_estimators=700;; score=0.768 total time= 2.6s
[CV 4/5] END max_depth=30, max_features=auto, min_samples_leaf=2, min_samples_split=100, n_estimators=700;; score=0.823 total time= 2.6s
[CV 5/5] END max_depth=30, max_features=auto, min_samples_leaf=2, min_samples_split=100, n_estimators=700;; score=0.847 total time=

2.6s

```
[CV 1/5] END max_depth=15, max_features=auto, min_samples_leaf=5,
min_samples_split=5, n_estimators=800;; score=0.878 total time= 3.9s
[CV 2/5] END max_depth=15, max_features=auto, min_samples_leaf=5,
min_samples_split=5, n_estimators=800;; score=0.888 total time= 3.8s
[CV 3/5] END max_depth=15, max_features=auto, min_samples_leaf=5,
min_samples_split=5, n_estimators=800;; score=0.778 total time= 3.8s
[CV 4/5] END max_depth=15, max_features=auto, min_samples_leaf=5,
min_samples_split=5, n_estimators=800;; score=0.841 total time= 3.9s
[CV 5/5] END max_depth=15, max_features=auto, min_samples_leaf=5,
min_samples_split=5, n_estimators=800;; score=0.853 total time= 4.0s
[CV 1/5] END max_depth=30, max_features=auto, min_samples_leaf=1,
min_samples_split=2, n_estimators=900;; score=0.826 total time= 6.6s
[CV 2/5] END max_depth=30, max_features=auto, min_samples_leaf=1,
min_samples_split=2, n_estimators=900;; score=0.863 total time= 6.6s
[CV 3/5] END max_depth=30, max_features=auto, min_samples_leaf=1,
min_samples_split=2, n_estimators=900;; score=0.748 total time= 6.6s
[CV 4/5] END max_depth=30, max_features=auto, min_samples_leaf=1,
min_samples_split=2, n_estimators=900;; score=0.835 total time= 6.7s
[CV 5/5] END max_depth=30, max_features=auto, min_samples_leaf=1,
min_samples_split=2, n_estimators=900;; score=0.824 total time= 6.6s
[CV 1/5] END max_depth=10, max_features=auto, min_samples_leaf=10,
min_samples_split=5, n_estimators=300;; score=0.880 total time= 1.2s
[CV 2/5] END max_depth=10, max_features=auto, min_samples_leaf=10,
min_samples_split=5, n_estimators=300;; score=0.893 total time= 1.2s
[CV 3/5] END max_depth=10, max_features=auto, min_samples_leaf=10,
min_samples_split=5, n_estimators=300;; score=0.782 total time= 1.3s
[CV 4/5] END max_depth=10, max_features=auto, min_samples_leaf=10,
min_samples_split=5, n_estimators=300;; score=0.845 total time= 1.2s
[CV 5/5] END max_depth=10, max_features=auto, min_samples_leaf=10,
min_samples_split=5, n_estimators=300;; score=0.858 total time= 1.3s
[CV 1/5] END max_depth=5, max_features=log2, min_samples_leaf=10,
min_samples_split=5, n_estimators=200;; score=0.847 total time= 0.6s
[CV 2/5] END max_depth=5, max_features=log2, min_samples_leaf=10,
min_samples_split=5, n_estimators=200;; score=0.873 total time= 0.6s
[CV 3/5] END max_depth=5, max_features=log2, min_samples_leaf=10,
min_samples_split=5, n_estimators=200;; score=0.755 total time= 0.6s
[CV 4/5] END max_depth=5, max_features=log2, min_samples_leaf=10,
min_samples_split=5, n_estimators=200;; score=0.817 total time= 0.6s
[CV 5/5] END max_depth=5, max_features=log2, min_samples_leaf=10,
min_samples_split=5, n_estimators=200;; score=0.820 total time= 0.6s
[CV 1/5] END max_depth=30, max_features=log2, min_samples_leaf=1,
min_samples_split=5, n_estimators=400;; score=0.852 total time= 1.8s
[CV 2/5] END max_depth=30, max_features=log2, min_samples_leaf=1,
min_samples_split=5, n_estimators=400;; score=0.885 total time= 1.7s
[CV 3/5] END max_depth=30, max_features=log2, min_samples_leaf=1,
min_samples_split=5, n_estimators=400;; score=0.763 total time= 1.7s
[CV 4/5] END max_depth=30, max_features=log2, min_samples_leaf=1,
min_samples_split=5, n_estimators=400;; score=0.831 total time= 1.7s
```

```
[CV 5/5] END max_depth=30, max_features=log2, min_samples_leaf=1,  
min_samples_split=5, n_estimators=400;; score=0.840 total time= 1.7s
```

```
RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(),  
                    param_distributions={'max_depth': [5, 10, 15, 20,  
25, 30],  
                                         'max_features': ['auto',  
'log2'],  
                                         'min_samples_leaf': [1, 2, 5,  
10],  
                                         'min_samples_split': [2, 5,  
10, 15,  
100],  
                                         'n_estimators': [100, 200,  
300, 400,  
500, 600,  
700, 800,  
900, 1000,  
1100,  
1200]}},  
                    random_state=2, verbose=5)
```

```
hyper_tuning.best_params_
```

```
{'n_estimators': 300,  
'min_samples_split': 5,  
'min_samples_leaf': 10,  
'max_features': 'auto',  
'max_depth': 10}
```

```
RF_hyper=RandomForestRegressor(n_estimators=700,min_samples_split=5,min  
n_samples_leaf=10,max_features='auto',max_depth=20)
```

```
RF_hyper.fit(X_train_transformed,y_train)
```

```
RandomForestRegressor(max_depth=20, max_features='auto',  
min_samples_leaf=10,  
                    min_samples_split=5, n_estimators=700)
```

```
tuning_predict=RF_hyper.predict(X_test_transformed)
```

```
print('Random Forest Regressor hyper parametric tuning for  
training :',RF_hyper.score(X_train_transformed,y_train))
```

```
Random Forest Regressor hyper parametric tuning for training :  
0.8841357926887093
```

```
print('Random Forest Regressor hyper parametric tuning for  
testing :',RF_hyper.score(X_test_transformed,y_test))
```

```
Random Forest Regressor hyper parametric tuning for testing :  
0.8773774832826662
```

GB Boosting Regressor

```
from sklearn.ensemble import GradientBoostingRegressor
GB_regressor = GradientBoostingRegressor()
GB_regressor.fit(X_train_transformed, y_train)

GradientBoostingRegressor()

y_test_pred = GB_regressor.predict(X_test_transformed)

#Testing the model accuracy

r2 = r2_score(y_test, y_test_pred) # Getting the R-squared (R²)
mae = mean_absolute_error(y_test, y_test_pred) # Getting the Mean Absolute Error
mse = mean_squared_error(y_test, y_test_pred) # Getting the Mean Squared Error
rmse = np.sqrt(mse) # Getting the Root Mean Squared Error

#Defining a DataFrame
GB_results = pd.DataFrame([['GBDT Regression', r2*100, mae, mse, rmse]],
                           columns=['Model', 'R-squared (R²)', 'Mean Absolute Error', 'Mean Squared Error', 'Root Mean Squared Error'])

#Printing the Model Results with different Metrics
results = results.append(GB_results, ignore_index=True)
results
```

	Model	R-squared (R²)	Mean Absolute Error \
0	Linear Regression	75.40	3938.02
1	KNN Regression	79.97	3006.35
2	Decison-tree Regression	58.68	3648.23
3	Random Forest Regression	82.55	2701.47
4	GBDT Regression	87.66	2310.81

	Mean Squared Error	Root Mean Squared Error
0	31503873.50	5612.83
1	25660048.66	5065.57
2	52924035.93	7274.89
3	22344392.41	4726.99
4	15805696.85	3975.64

```
print('GBDT Regression Train Score is : ',
      GB_regressor.score(X_train_transformed, y_train))
print('GBDT Regression Test Score is : ',
      GB_regressor.score(X_test_transformed, y_test))
```

```
GBDT Regression Train Score is : 0.8993791860213253
GBDT Regression Test Score is : 0.8765956618967508
```

Residuals (Difference between y_test and y_pred)

```
test_residual= y_test - y_test_pred

medical_GB = pd.DataFrame({'Actual': y_test, 'Predicted': y_test_pred,
'diff':(abs(y_test-y_test_pred))})

medical_GB.head()
```

	Actual	Predicted	diff
12	1826.84	4224.89	2398.05
306	20177.67	6809.64	13368.03
318	7421.19	9797.54	2376.35
816	2842.76	5786.27	2943.51
157	15518.18	15694.14	175.96

```
results[["Model", "R-squared (R²)"]]
```

	Model	R-squared (R²)
0	Linear Regression	75.40
1	KNN Regression	79.97
2	Decison-tree Regression	58.68
3	Random Forest Regression	82.55
4	GBDT Regression	87.66

| Conclusion

In conclusion, this machine learning project successfully developed a predictive model for medical cost estimation based on the provided dataset.

The GBDT regression model showed strong performance in predicting medical costs, taking into account various features such as age, BMI, smoking status, and more.

This model can be utilized to assist in estimating medical costs for individuals and provide insights for healthcare providers and insurance companies.